

Pawan Mehta

✉ pawanm571@gmail.com | 📞 +917004581187
🌐 <https://www.linkedin.com/in/pawanm571/> | 📍 Hyderabad, India

Profile

Results-driven Data Engineer with over 5 years of experience designing and managing scalable data pipelines and warehouses using Snowflake, AWS, and GCP. Expert in ETL/ELT processes, data modeling, and data governance, with proficiency in Python, SQL, Apache Airflow, and DBT. Skilled in collaborating with data scientists, product managers, and business stakeholders to deliver production-ready data solutions aligned with business roadmaps. Adept at Agile methodologies, quality assurance, and resolving technical issues to ensure data quality and operational efficiency. Holds AWS certifications and a Bachelor's degree in Computer Applications, with a strong focus on delivering high-impact data infrastructure.

Skills

- **Data Engineering:** Data Pipelines, ETL/ELT, Data Warehousing, Data Modeling
- **Programming & Tools:** Python, SQL, Apache Spark, Apache Airflow, DBT, Great Expectations
- **Cloud Platforms:** Snowflake, AWS, GCP
- **Infrastructure Automation:** Terraform, Jenkins, Git, Kubernetes, Docker
- **Data Governance:** Data Quality, Compliance, Snowflake Access Controls, AWS KMS
- **Visualization:** PowerBI

Work Experience

Mouri Tech, Hyderabad

Oct 2024 - Present

Senior Data Engineer

- Designed and launched scalable data pipelines in Snowflake, processing 100TB+ daily IoT and supply chain data using Snowpipe and DBT, achieving 95% operational efficiency for 1M+ transactions.
- Developed star schema data models in Snowflake with clustering keys and partitioning, collaborating with product managers to reduce reporting latency by 20% for real-time analytics.
- Implemented enterprise-grade data governance frameworks using Great Expectations, ensuring data quality and compliance for 1M+ customer records across business domains.
- Automated ETL/ELT workflows with Python and Airflow, integrating with AWS S3 and Lambda, reducing data processing time by 25% and enabling seamless data management.
- Led Agile ceremonies (standups, sprint planning, retrospectives), aligning data solutions with product roadmaps and stakeholder requirements, cutting project delays by 15%.
- Spearheaded quality assurance by defining Python-based testing strategies, identifying risks, and resolving technical issues, improving data reliability by 25%.
- Deployed Snowflake infrastructure on AWS using Terraform and Kubernetes, optimizing CI/CD pipelines with Jenkins and Git, reducing deployment time by 50%.
- Built strong relationships with business application owners, delivering Looker dashboards integrated with Snowflake for real-time insights, achieving 99.9% system uptime.

HCL TECH, Lucknow

Sept 2020 - sep 2024

Data Engineer

- Built robust Snowflake pipelines for 5TB+ daily SIEM logs using Snowpipe and DBT, enabling real-time analytics and reducing threat detection time to 15 minutes.
- Designed and optimized data warehouse schemas in Snowflake with SCD Type 2 and clustering keys, achieving 99.8% dashboard reliability for stakeholder reporting.
- Collaborated with business stakeholders and data scientists to gather requirements, delivering Tableau dashboards integrated with Snowflake for actionable security insights.
- Automated ETL/ELT processes with Python, Airflow, and DBT, migrating 3.7M+ legacy records to Snowflake, improving query performance by 35%.
- Implemented data quality frameworks with Great Expectations, ensuring compliance and traceability for 1M+ customer records across enterprise systems.
- Actively participated in Agile sprint reviews and retrospectives, maturing requirements gathering processes to align with business priorities, reducing delivery time by 20%.
- Optimized Snowflake queries with partitioning and indexing, troubleshooting performance issues to reduce latency by 30% for high-volume analytics.
- Managed AWS infrastructure with Terraform and Jenkins, deploying Kubernetes clusters for data applications, cutting setup time by 30%.

Education

Master of Computer Applications	University of Mysore, India	Jan 2025 - Present
Bachelor of Computer Applications	Amity University, Noida, India	June 2021 - June 2024

Projects

Sales DataMart

- Architected a Snowflake-based DataMart integrating 7+ ERP systems using Snowpipe and DBT, achieving 90% reporting efficiency for 1B+ transactions.
- Optimized star schemas and clustering keys, reducing query latency by 25% for real-time stakeholder analytics.

OTIF DataMart

- Developed a modular Snowflake DataMart with Snowpipe and Airflow, reducing supply chain delays by 15% through optimized ETL processes.
- Designed schemas for real-time analytics, supporting data-driven stakeholder decisions.

Financial DataMart

- Built a Snowflake-based Financial DataMart with Python and Snowpipe, integrating multi-ERP data, achieving 90% accuracy in daily balance reporting.
- Developed validation scripts to ensure data integrity, reducing errors by 10%.

History vs Forecast DataMart

- Created a unified Snowflake model for historical and forecast data, using Snowpark for ML and achieving 90% accuracy in variance analysis.
- Designed experiments to measure model performance, supporting stakeholder decision-making with actionable insights.

Certifications

- AWS Certified Developer – Associate
- DBT Fundamentals
- AWS Cloud Practitioner
- IBM Data Science Orientation

Languages

- Hindi (Native)
- English (Fluent)